


What is MLOps ?

How AI Success Demanded Operational Excellence

School of DevOps™

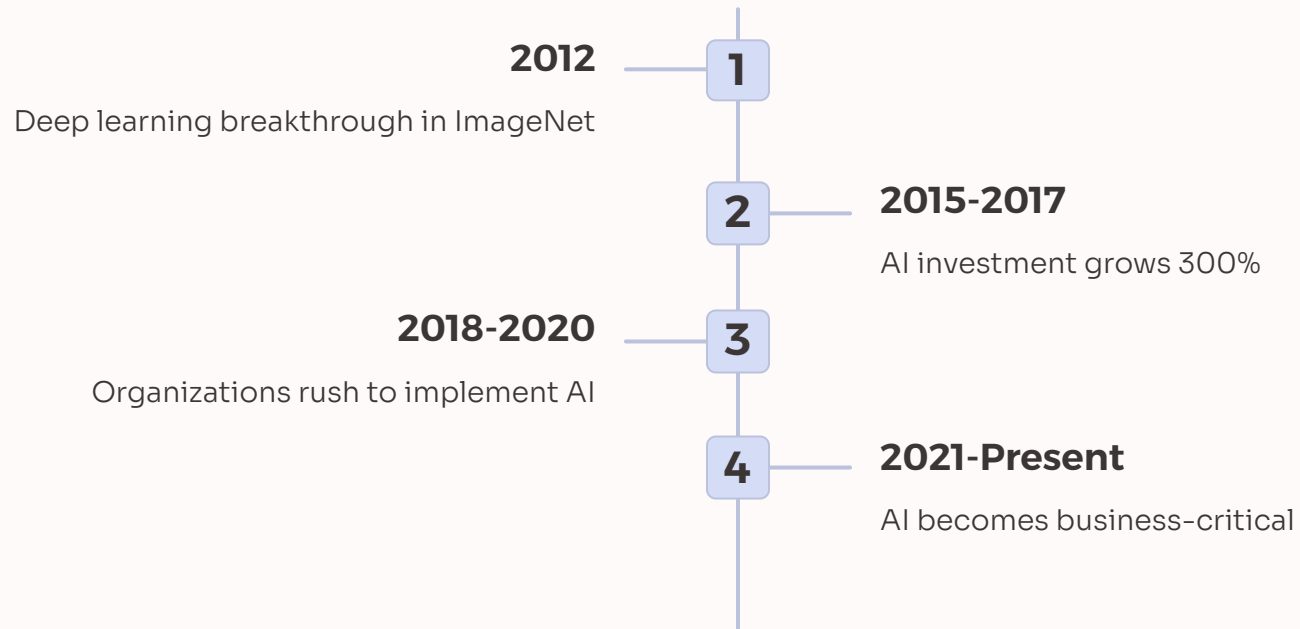


by **Gourav Shah**





The AI Revolution Begins



Great Expectations vs. Reality

The Dream

- Train a model
- Deploy it
- Watch magic happen
- Profit!

The Reality

- 87% of ML projects never reach production
- 9+ months from model to deployment
- 50%+ of models fail to deliver value



The Cold, Hard Truth

\$15.7T

AI Impact

Projected global economic impact by 2030

83%

Frustrated

Data scientists facing deployment challenges

90%

Struggling

Organizations having AI implementation issues

70%

Improvement

Reduction in time-to-value with proper practices

The 3AM Crisis

The Incident

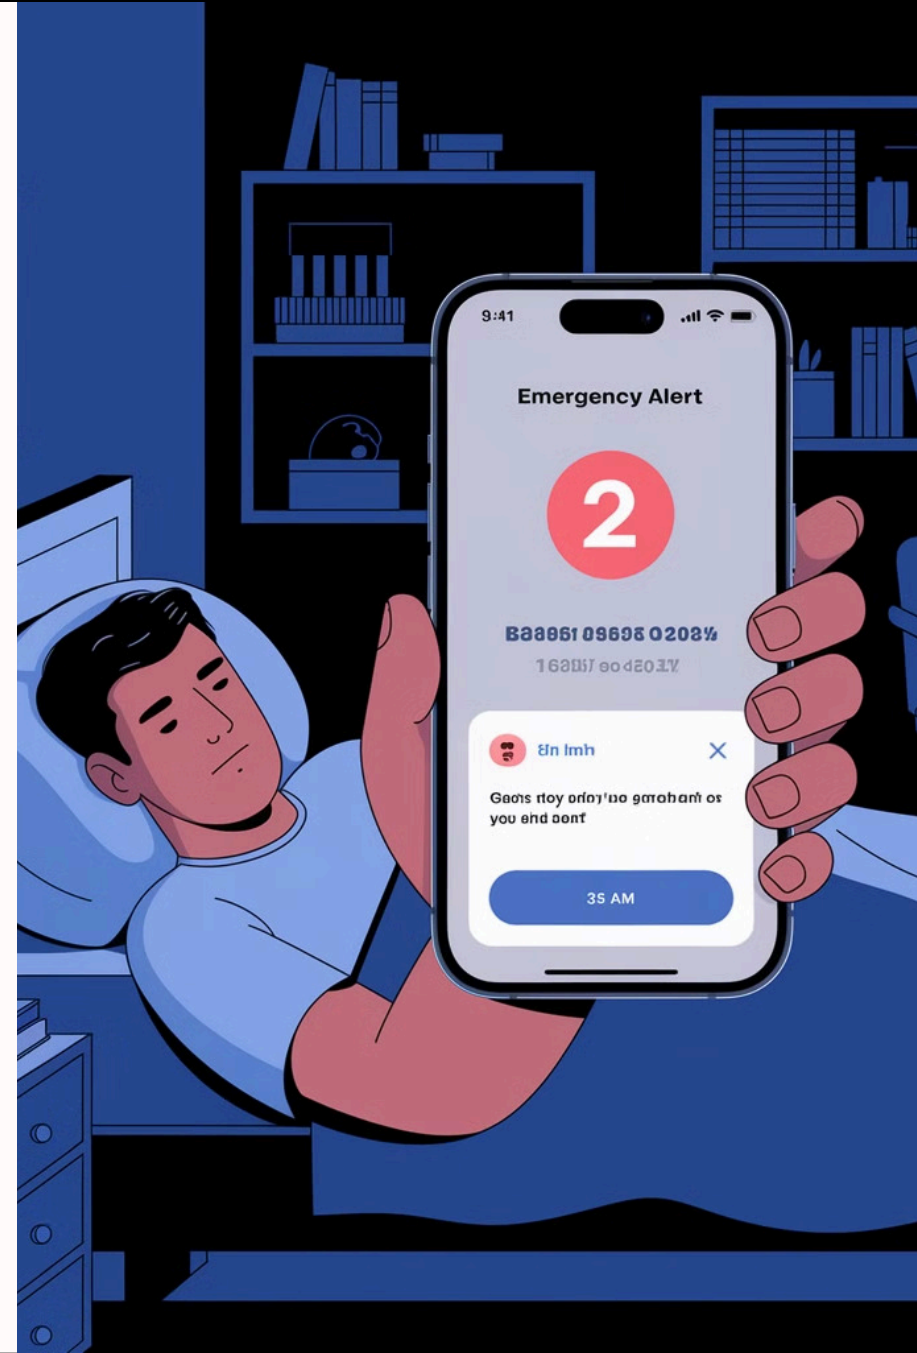
Recommendation engine suggesting winter coats to Australian users in summer

The Questions

Which model version? What training data? How did it pass testing?

The Problem

No systematic operational practices means no good answers



The Production Gap


Data Science World

- Jupyter notebooks
- Experimentation focus
- Static datasets
- Academic metrics


Production World

- Scalable infrastructure
- Reliability requirements
- Dynamic data
- Business metrics

The Hidden Complexity of ML Systems



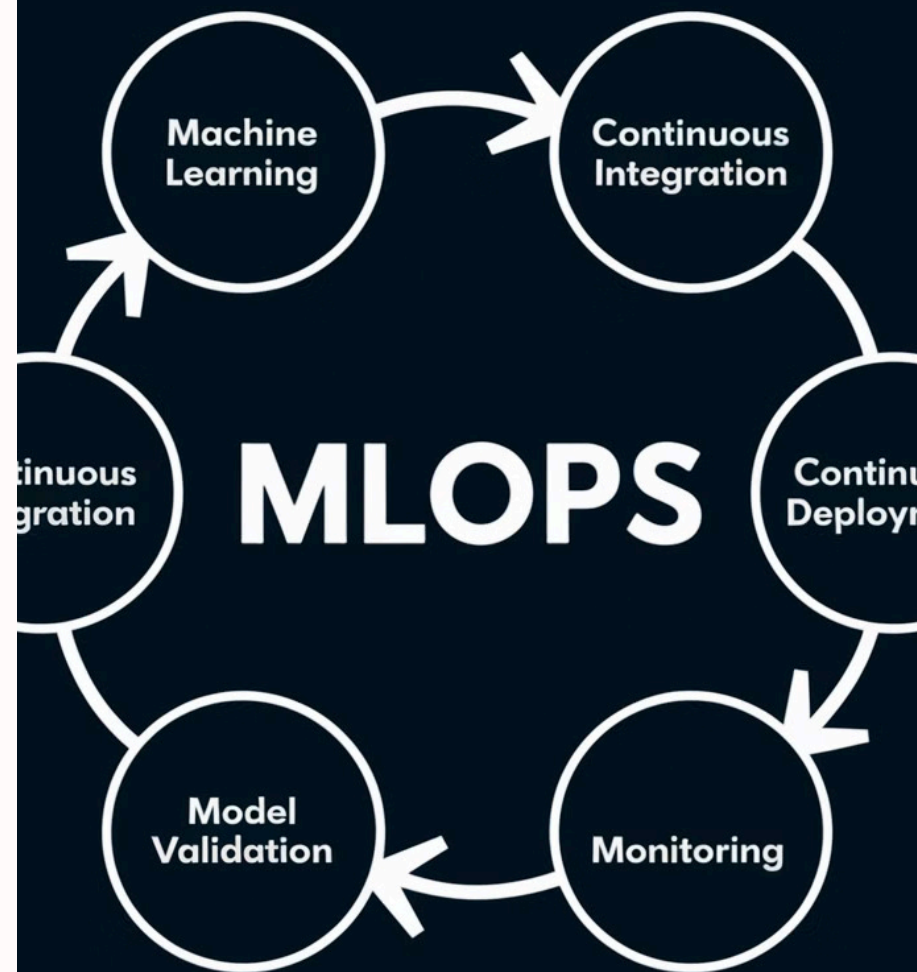
MLOps: A New Discipline Emerges



What is MLOps?

MLOps is a set of practices at the intersection of Machine Learning, DevOps, and Data Engineering aimed at deploying and maintaining ML systems in production reliably and efficiently.

- 1 Bridges development and operations
- 2 Standardizes the ML lifecycle
- 3 Automates repetitive processes
- 4 Enables reproducibility and governance



If ML Were a Restaurant...

Without MLOps

- No standardized recipes
- No ingredient tracking
- Inconsistent meals
- Can't scale successful dishes

With MLOps

- Recipe versioning
- Ingredient quality control
- Consistent preparation
- Scalable kitchen operations



The 3 Pillars of MLOps



CI/CD

Automated testing,
building, deployment



Orchestration & Automation


End-to-end
workflow
management



Monitoring & Management

Performance
tracking, drift
detection

RELIABLE ML IN PRODUCTION




MLOps Core Practices

Version Everything

Code, data, models, configs

Enable Governance

Lineage, documentation,
compliance



Automate Pipelines

Training, testing, deployment

Track Experiments

Parameters, metrics, artifacts

Monitor Continuously

Performance, drift, resources

The Technical Debt Monster

Machine learning systems have a special capacity for incurring technical debt.

1 ML-Specific Debt


Data dependencies,
configuration complexity

2 Experimentation Issues

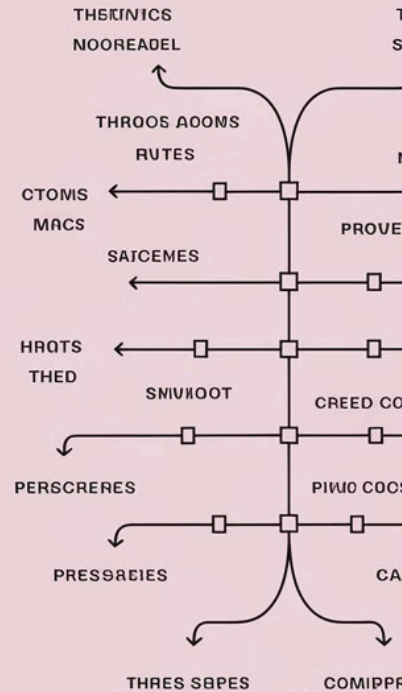
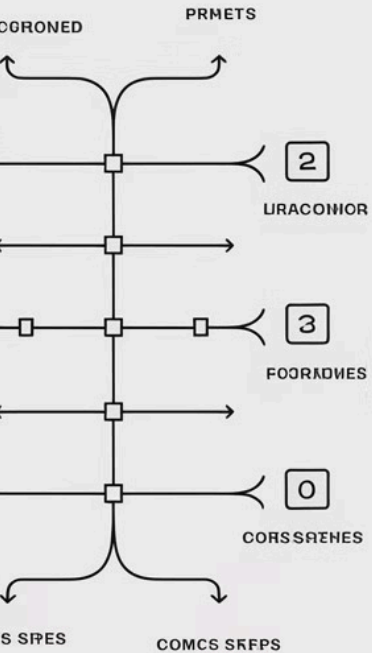
No tracking,
undocumented features

3 Operational Problems

Manual deployment, lack of monitoring



ML Lifecycle vs. Software Development



1

Traditional Software

Requirements → Design → Implementation → Testing → Deployment → Maintenance

2

ML Development

Problem framing → Data prep → Feature engineering → Training → Evaluation → Deployment → Monitoring

3

Key Differences

Data dependency, non-deterministic behavior, continuous retraining



The Bottom Line: Business Value

70%

Faster

Reduction in time-to-deployment

40%

Better

Improvement in model performance

65%

Reliable

Fewer production incidents

4x

Scalable

More models in production

The MLOps Maturity Journey



Level 0: Manual Process

Manual preparation, no versioning



Level 1: Pipeline Automation

Automated training, basic versioning



Level 2: CI/CD Automation

Automated testing, inference deployment, basic monitoring



Level 3: Automated Operations

Drift detection, on-demand retraining



Level 4: Full Automation

Auto-triggered retraining, self-healing

The Pioneers' Advantage



Netflix

Created Metaflow, reduced deployment time by 60%



Uber

Built Michelangelo, enabling millions of daily predictions




Facebook

Developed FBLearn, supporting 1M+ model runs daily



Airbnb

Implemented Bighead, increased experiment velocity 4x



The Evolving AI Landscape

1

2015-2018: Traditional ML Focus

Custom models, structured data, centralized development

2

2018-2021: Deep Learning Expansion

Neural networks, unstructured data, model ensembles

3

2021-Present: Foundation Models & Agents

Large language models, multimodal systems, agentic capabilities



The Evolution Continues

1

MLOps

Traditional machine learning operations

2

LLMOps

Foundation model operations

3

AgenticAIOps

Autonomous agent operations

Reflections

1 Maturity Assessment

Where are you on the MLOps journey?

2 Pain Points

Biggest challenge moving ML to production?

3 Time Savings

How much time could proper MLOps save?

4 Business Impact

Value of deploying models twice as fast?

